



Computación Concurrente, Paralela y Distribuida

Leandro Rodríguez Liñares – David Olivieri
Curso 2024/25



Práctica 1: Introducción



Descargar e instalar Anaconda

- Descargar anaconda:

<https://www.anaconda.com/download>

(no es necesario registrarse)

- Ejecutar el instalador:

```
$ chmod +x Anaconda3-2024.10-1-Linux-x86_64.sh  
$ ./Anaconda3-2024.10-1-Linux-x86_64.sh
```

Importante: contestar “yes” al final de la instalación

El camino en que se instala será: `/home/my_user/anaconda3` (6.5GB)

- ▶ Crear y activar un nuevo entorno en Anaconda

```
$ conda create --name mpi  
$ conda activate mpi
```

- ▶ Instalar MPI:

```
(mpi) $ conda install ipython jupyter numpy mpi4py openmpi  
(mpi) $ conda install -c conda-forge gcc
```

- ▶ Comprobar funcionamiento:

```
#!/usr/bin/env python  
  
from mpi4py import MPI  
  
comm = MPI.COMM_WORLD  
my_rank = comm.rank  
num_processes = comm.size  
  
if my_rank != 0:  
    data = "Saludos del proceso {}".format(my_rank)  
    comm.send(data, dest=0)  
  
else:  
    print("Hola, soy el proceso %d/%d y recibo:" % (my_rank, num_processes))  
    for source_rank in range(1, num_processes):  
        data_in = comm.recv(source = source_rank)  
        print("    "+data_in)
```

```
(mpi) $ chmod +x CheckMPI.py  
$ mpirun -np 3 ./CheckMPI.py  
Hola, soy el proceso 0/3 y recibo:  
    Saludos del proceso 1  
    Saludos del proceso 2  
$ mpirun --oversubscribe -np 6 ./CheckMPI.py  
Hola, soy el proceso 0/6 y recibo:  
    Saludos del proceso 1  
    Saludos del proceso 2  
    Saludos del proceso 3  
    Saludos del proceso 4  
    Saludos del proceso 5
```

```

#include <stdio.h>
#include <string.h>
#include "mpi.h"
int main(int argc, char* argv[])
{
    int my_rank;
    int p;
    int source; int dest;
    int tag = 0;
    char message[100];

    MPI_Status status;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &p);

    if (my_rank != 0) {
        sprintf(message, "Saludos del proceso %d", my_rank);
        dest = 0;
        MPI_Send(message, strlen(message) + 1, MPI_CHAR, dest, tag, MPI_COMM_WORLD);
    }
    else {
        printf("Hola, soy el proceso %d (hay %d procesos) y recibo:\n", my_rank, p);
        for (source = 1; source < p; source++) {
            MPI_Recv(message, 100, MPI_CHAR, source, tag, MPI_COMM_WORLD, &status);
            printf("%s\n", message);
        }
    }

    MPI_Finalize();

    return 0;
}

```

```

(mpi) $ mpicc CheckMPI.c -o CheckMPI
$ mpirun -np 4 ./CheckMPI
Hola, soy el proceso 0 (hay 4 procesos) y recibo:
Saludos del proceso 1
Saludos del proceso 2
Saludos del proceso 3

```

Instalación Anaconda: 11GB

Instalar Numba, Cython, NVidia CUDA y Tensorflow

- ▶ Crear y activar un nuevo entorno en Anaconda

```
$ conda create --name gpu  
$ conda activate gpu
```

- ▶ Instalar NVidia CUDA y herramientas:

```
(gpu) $ conda install numba scipy matplotlib  
cudatoolkit ipython jupyter cython
```

- ▶ Comprobar funcionamiento de GPU:

```
(gpu) $ python3 CheckGPU.py  
Found 1 device(s).  
Device: 0  
  Name: NVIDIA T400 4GB  
  Compute Capability: 7.5  
  Multiprocessors: 6  
  CUDA Cores: 384  
  Concurrent threads: 6144  
  GPU clock: 1425 MHz  
  Memory clock: 5001 MHz  
  Total Memory: 3901 MiB  
  Free Memory: 3025 MiB
```



Python vs Numba vs Cython



```
# Cython Function
import pyximport; pyximport.install()
from sum_series import sum_series_cython
```

```
#Python Function
def sum_series_python(x):
    y = 0
    for i in range(x):
        y += i
    return y
```

```
# Numba Function
from numba import njit
@njit(cache = True)
def sum_series_numba(x):
    y = 0
    for i in range(x):
        y += i
    return y
```

PythonNumbaCython.py

```
# Cython Function
def sum_series_cython(int x):
    cdef int y = 0
    cdef int i
    for i in range(x):
        y += i
    return y
```

sum_series.pyx

```
(gpu) leandro@T1700:~/Tmp$ ipython
```

```
...
```

```
In [1]: %run PythonNumbaCython.py
```

```
...
```

```
In [2]: %timeit sum_series_python(1000000)
```

```
32.1 ms ± 433 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)
```

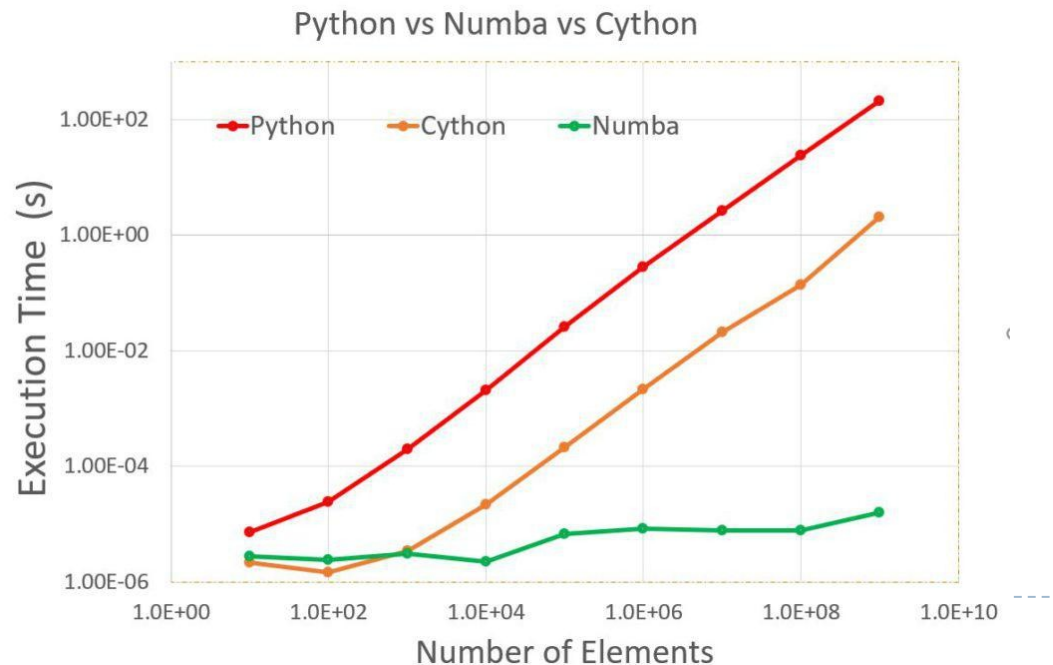
```
...
```

1000 elementos

- python 25.2 µs
- cython 301 ns (S = 83.7)
- numba 146 ns (S=172.6)

1000000 elementos

- python 31.8 ms
- cython 261 µs (S = 121.8)
- numba 151 ns (S=210596)





```
from numba import cuda
import numpy as np

blocks_per_grid = 1000
threads_per_block = 1024
n = threads_per_block*blocks_per_grid

def add_cpu(x, y, out):
    out = x + y;

@cuda.jit
def add_gpu(x, y, out):
    idx = cuda.grid(1)
    out[idx] = x[idx] + y[idx]

h_x = np.ones(n) # [1...1]
h_y = np.ones_like(h_x)
h_out = np.ones_like(h_x)

d_x = cuda.to_device(h_x)
d_y = cuda.to_device(h_y)
d_out = cuda.device_array_like(d_x)

add_cpu(h_x, h_y, h_out)

# Un hilo para cada elemento
add_gpu[blocks_per_grid, threads_per_block](d_x, d_y, d_out)

print("Número de hilos:", n)
print(d_out.copy_to_host()) # Resultado: [2...2]
```

```
(gpu) $ ipython
...
In [1]: %run Cuda_add.py
Número de hilos: 1024000
[2. 2. 2. ... 2. 2. 2.]
In [2]: %timeit add_cpu(h_x, h_y, h_out)
1.69 ms ± 601 ns per loop (mean ± std. dev. of 7
runs, 1,000 loops each)
In [3]: %timeit add_gpu[blocks_per_grid,
threads_per_block](d_x, d_y, d_out)
359 µs ± 138 ns per loop (mean ± std. dev. of 7 runs,
10,000 loops each)
```

- 1 bloque (1024 hilos):
S = 0.026
- 100 bloques (102400 hilos):
S = 2.025
- 10000 bloques (10240000 hilos):
S = 5.496

Instalar Tensorflow



► Instalar Tensorflow:

```
(gpu) $ python3 -m pip install tensorflow[and-cuda]
```

```
....
```

► Probar instalación:

```
(gpu) $ ipython
```

```
...
```

```
In [1]: import tensorflow as tf
```

```
...
```

```
In [2]: print("Num GPUs available: ",  
len(tf.config.list_physical_devices('GPU')))
```


```
...
```

```
Num GPUs available: 1
```

A screenshot of a terminal window showing the output of the command 'python3 -m pip install tensorflow[and-cuda]'. The terminal displays a list of packages being collected and downloaded, including flatbuffers, gast, h5py, absl-py, grpcio, astunparse, ml-dtypes, typing_extensions, nvidia_cublas, nvidia_cufft, and nvidia_cuda. The progress bars indicate the download status for each package.

Instalación Anaconda: 19GB

Probar Tensorflow



```
(gpu) $ python3 ./CheckTF.py
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	401920
dense_1 (Dense)	(None, 512)	262656
dense_2 (Dense)	(None, 10)	5130

Total params: 669706 (2.55 MB)

Trainable params: 669706 (2.55 MB)

Non-trainable params: 0 (0.00 Byte)

Epoch 1/10

1875/1875 [==....==] - 5s 2ms/step - loss: 0.1874 - accuracy: 0.9436 - val_loss: 0.1054 - val_accuracy: 0.9683

Epoch 2/10

1875/1875 [==....==] - 4s 2ms/step - loss: 0.0869 - accuracy: 0.9751 - val_loss: 0.1311 - val_accuracy: 0.9647

Epoch 3/10

1875/1875 [==....==] - 4s 2ms/step - loss: 0.0606 - accuracy: 0.9829 - val_loss: 0.0810 - val_accuracy: 0.9780

Epoch 4/10

1875/1875 [==....==] - 4s 2ms/step - loss: 0.0466 - accuracy: 0.9870 - val_loss: 0.0959 - val_accuracy: 0.9777

Epoch 5/10

1875/1875 [==....==] - 4s 2ms/step - loss: 0.0380 - accuracy: 0.9895 - val_loss: 0.0718 - val_accuracy: 0.9842

Epoch 6/10

1875/1875 [==....==] - 4s 2ms/step - loss: 0.0295 - accuracy: 0.9923 - val_loss: 0.0963 - val_accuracy: 0.9811

Epoch 7/10

1875/1875 [==....==] - 4s 2ms/step - loss: 0.0230 - accuracy: 0.9936 - val_loss: 0.1180 - val_accuracy: 0.9801

Epoch 8/10

1875/1875 [==....==] - 4s 2ms/step - loss: 0.0199 - accuracy: 0.9946 - val_loss: 0.1119 - val_accuracy: 0.9812

Epoch 9/10

1875/1875 [==....==] - 4s 2ms/step - loss: 0.0148 - accuracy: 0.9959 - val_loss: 0.1149 - val_accuracy: 0.9820

Epoch 10/10

1875/1875 [==....==] - 4s 2ms/step - loss: 0.0132 - accuracy: 0.9963 - val_loss: 0.1293 - val_accuracy: 0.9810



► Instalar PyTorch:

```
(gpu) $ conda install pytorch torchvision torchaudio pytorch-cuda -c pytorch -c nvidia
....
```

► Probar instalación:

```
(gpu) $ ipython
...
In [1]: import torch
...
In [2]: print(torch.__version__)
2.5.1

In [3]: torch.cuda.is_available()
Out[3]: True

In [4]: torch.cuda.device_count()
Out[4]: 1

In [5]: torch.cuda.get_device_name(0)
Out[5]: 'NVIDIA T400 4GB'
```

Probar Pytorch

```
(gpu) $ python3 ./CheckTF.py
...
Extracting /tmp/data/MNIST/raw/t10k-labels-idx1-ubyte.gz to /tmp/data/MNIST/raw

Length of dataset: 60000
Length of first vector in dataset: torch.Size([784])
Label of first vector in dataset: 5
Length of Train Dataset: 50000
Length of Test Dataset: 10000
Sequential(
  (0): Linear(in_features=784, out_features=512, bias=True)
  (1): ReLU()
  (2): Linear(in_features=512, out_features=10, bias=True)
  (3): Softmax(dim=1)
)
Epoch 01: loss 1.5777 - accuracy 0.8966 - validation accuracy 0.9294
Epoch 02: loss 1.5177 - accuracy 0.9496 - validation accuracy 0.9497
Epoch 03: loss 1.5033 - accuracy 0.9623 - validation accuracy 0.9548
Epoch 04: loss 1.4942 - accuracy 0.9702 - validation accuracy 0.9672
Epoch 05: loss 1.4884 - accuracy 0.9761 - validation accuracy 0.9626
Epoch 06: loss 1.4841 - accuracy 0.9798 - validation accuracy 0.9728
Epoch 07: loss 1.4809 - accuracy 0.9825 - validation accuracy 0.9698
Epoch 08: loss 1.4786 - accuracy 0.9844 - validation accuracy 0.9746
Epoch 09: loss 1.4763 - accuracy 0.9865 - validation accuracy 0.9699
Epoch 10: loss 1.4747 - accuracy 0.9880 - validation accuracy 0.9764
```

Instalación Anaconda: 27GB

Borrar anaconda

```
(gpu) $ conda activate base
(base) $ conda init --reverse -all
...
==> For changes to take effect, close and re-open your current shell. <==

$ rm -rf ~/anaconda3
$ rm -rf ~/.condarc ~/.conda
```